



计算机视觉表征与识别

Chapter 5: Edges

王利民

媒体计算课题组

<http://mcg.nju.edu.cn/>



Image pyramids



Gaussian




Progressively blurred and subsampled versions of the image. Adds scale invariance to fixed-size algorithms.

Laplacian



Shows the information added in Gaussian pyramid at each spatial scale. Useful for noise reduction & coding.

Steerable pyramid



Shows components at each scale and orientation separately. Non-aliased subbands. Good for texture and feature analysis.



Linear Image Transforms



1297x256










Image representation




- Pixels: great for spatial resolution, poor access to frequency
- Fourier transform: great for frequency, not for spatial information
- Pyramids/filter banks: balance between spatial and frequency information


Vision: a multi-stage network



Model of Retinal and LGN cells






Difference of Gaussian Model



- The output of a retinal ganglion cell is a weighted sum of its inputs. You can treat the cell as a shift-invariant linear system
- **Difference of Gaussian (DoG):** A model of the operation performed by retinal and ganglion cells (LGN). The DoG model supposes that the neural response results from the combined signal of two separate mechanisms
- LGN and retinal neurons have circular receptive fields: **they respond equally well to all stimulus orientations, and at different spatial frequency (scale)**
- Multiple DoG is called a Laplacian Pyramid

Oriented representation in V1

- Simple-cell receptive fields (RFs) are constructed from the output of LGN cells
- They are selective to oriented contours and edges



RECEPTIVE FIELD STRUCTURE

GABOR PATCH




What is texture?



- **Texture is a phenomenon that is widespread, easy to recognize, and hard to define.**
- Views of large numbers of small objects
- Regular or stochastic patterns caused by bumps, grooves, and/or markings
- Textures tend to show **repetition**: the same local patch appears again and again.



Texture overview





Why analyze texture?



Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.





Texture-related tasks



- **Shape from texture**
 - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
 - Analyze, represent texture
 - Group image regions with consistent texture
- **Synthesis**
 - Generate new texture patches/images given some examples



Analysis vs. Synthesis





Texture representation



- Textures are made up of repeated local patterns, so:
 - Find the patterns
 - Use filters that look like patterns (spots, bars, raw patches...)
 - Consider magnitude of response
 - Describe their statistics within each local window, e.g.,
 - Mean, standard deviation
 - Histogram
 - Histogram of “prototypical” feature occurrences




Filter banks



- Our previous example used two filters, and resulted in a 2-dimensional feature vector to describe texture in a window.
 - x and y derivatives revealed something about local structure.
- We can generalize to apply a collection of multiple (d) filters: a “filter bank”
- Then our feature vectors will be d -dimensional.
 - still can think of nearness, farness in feature space



Filter banks



- What filters to put in the bank?
 - Typically we want a combination of scales and orientations, different types of patterns.

Matlab code available for these examples:

<http://www.robots.ox.ac.uk/~vgg/research/texclass/filters.html>




2D Textons



- Goal: find canonical local features in a texture;

- 1) Filter image with linear filters:






- 2) Vector quantization (k-means) on filter outputs;
- 3) Quantization centers are the textons.

- Spatial distribution of textons defines the texture;



2D Textons (cont' d)





Texture/Material Representation



- Each material is now represented as a spatial arrangement of symbols from the texton vocabulary;
- Recognition: ignore spatial arrangement, use histogram ($K=100$);



Similarity of materials



- Similarity between histograms measured using chi-square difference:


$$\chi^2(h_1, h_2) = \sum_{n=1}^N \frac{(h_1(n) - h_2(n))^2}{h_1(n) + h_2(n)}$$



Texture synthesis



- Goal: create new samples of a given texture
- Many applications: virtual environments, hole-filling, texturing surfaces







The Challenge

- Need to model the whole spectrum: from repeated to stochastic texture

Alexei A. Efros and Thomas K. Leung, “Texture Synthesis by Non-parametric Sampling,” Proc. International Conference on Computer Vision (ICCV), 1999.




repeated




stochastic




Both? 22



Synthesizing One Pixel



input image




synthesized image


- What is $P(x|\text{neighborhood of pixels around } x)$?
- Find all the windows in the image that match the neighborhood
- To synthesize x
 - pick one matching window at random
 - assign x to be the center pixel of that window



Image Quilting [Efros & Freeman 2001]



non-parametric
sampling



Input image

Synthesizing a block

- Observation: neighbor pixels are highly correlated
- Idea: unit of synthesis = block

- Exactly the same but now we want $P(B|N(B))$
- Much faster: synthesize all pixels in a block at once




Today's class



- Introduction to edge detection.
- Gradients and edges.
- Canny edge detector.
- Object contour.
- Pb edge detector.
- Recent advances in edge detection.
- Straight line detection





From texture to edge/segmentation






From Pixels to Perception





**Mid-level operations of
Segmentation and Grouping**



Recognition, Reconstruction & Reorganization



Recognition



Reconstruction

Reorganization



Figure / Ground

Finding groups of pixels that go together
(parts, objects, textures, holes)



Predicted scene categories:

forest - broadleaf (0.498), swimming hole (0.402), bayou (0.062)

Predicted scene categories:

forest - broadleaf (0.979)



Figure / Ground





Figure / Ground









First problem: edge detection



- Goal: compute something like a line drawing of a scene





Why




- Edges reflect intrinsic properties of a scene
 - Capture shape information
 - Independent of illumination
- Our human visual system does something like this
- Good initial step for solving other problems
 - Recognition, tracking, etc.



Why



- Extract information, recognize objects
- Recover geometry and viewpoint





Issues



- No precise problem formulation
- Much harder than it seems to be
- **Edge detectors** usually work by detecting “big changes” in image intensity
- **Boundary** is contour in the image plane that represents a change in pixel ownership from object or surface to another.



Edge detection



- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.




Figure from J. Shotton et al., PAMI 2007


- **Main idea:** look for strong gradients, post-process



Recall: How much light is recorded



- Major factors
 - Illumination strength and direction
 - Surface geometry
 - Surface material
 - Nearby surfaces
 - Camera gain/exposure






What causes an edge?



Reflectance change:
appearance
information, texture

Change in surface
orientation: shape




Depth discontinuity:
object boundary

Cast shadows

Slide credit:
Kristen Grauman




Edges/gradients and invariance





Derivatives and edges

An edge is a place of rapid change in the image intensity function.





Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement above as convolution, what would be the associated filter?

Slide credit: Kristen Grauman



Partial derivatives of an image





$$\frac{\partial f(x, y)}{\partial x}$$



$$\frac{\partial f(x, y)}{\partial y}$$

-1	1
----	---



-1	?	1
1	or	-1

Which shows changes with respect to x?

(showing filters for correlation)



Finite difference filters



- Other approximations of derivative filters exist:

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

Source: K. Grauman






Image gradient



The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity




The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$





The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$





Effects of noise






Effects of noise




Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal

$$f(x)$$



$$\frac{d}{dx} f(x)$$



Where is the edge?



Effects of noise




- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?

Source: D. Forsyth




Solution: smooth first

f



h



Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$




Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$


Differentiation property of convolution.

Sigma = 50


f



$\frac{\partial}{\partial x}h$



$(\frac{\partial}{\partial x}h) \star f$






Derivative of Gaussian filters



$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$


$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

$$\otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$








Derivative of Gaussian filters



x-direction



y-direction





The Sobel Operator: A common approximation of derivative of gaussian



- Common approximation of derivative of Gaussian

$$\frac{1}{8} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

s_x

$$\frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

s_y

- The standard defn. of the Sobel operator omits the $1/8$ term
 - doesn't make a difference for edge detection
 - the $1/8$ term **is** needed to get the right gradient value



Mask properties



■ Smoothing

- Values positive
- Sum to 1 → constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components; “low-pass” filter

■ Derivatives

- _____ signs used to get high response in regions of high contrast
- Sum to ___ → no response in constant regions
- High absolute value at points of high contrast






Smoothing with a Gaussian



Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...






Effect of σ on derivatives



$\sigma = 1$ pixel





$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.
Larger values: larger scale edges detected
Smaller values: finer features detected



So, what scale to choose?

It depends what we're looking for.






Laplacian of Gaussian




Consider $\frac{\partial^2}{\partial x^2}(h \star f)$


f



$\frac{\partial^2}{\partial x^2} h$



$(\frac{\partial^2}{\partial x^2} h) \star f$



Where is the edge?


Zero-crossings of bottom graph

58

Slide credit: Steve S



2D edge detection filters



Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$\nabla^2 h_\sigma(u, v)$
Laplacian of Gaussian



- ∇^2 is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$


Slide credit: Steve Seitz



Designing an edge detector



- Criteria for an “optimal” edge detector:
 - **Good detection:** minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)
 - **Good localization:** must be as close as possible to the true edges
 - **Single response:** the detector must return one point only for each true edge point;




Evaluate Edge Detection

$$\text{precision} = \frac{\text{GT} \cap \text{RM}}{\text{RM}} = \frac{\text{TP}}{\text{RM}}$$

$$\text{recall} = \frac{\text{GT} \cap \text{RM}}{\text{GT}} = \frac{\text{TP}}{\text{GT}}$$

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



Ground Truth (GT)

Results of Method (RM)

True Positives (TP)

True Negatives (TN)

False Negatives (FN)

False Positives (FP)

Basic Comparisons of Edge Operators

Gradient:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Good Localization
Noise Sensitive
Poor Detection

Roberts (2 x 2):

0	1
-1	0

1	0
0	-1

Sobel (3 x 3):

-1	0	1
-1	0	1
-1	0	1

1	1	1
0	0	0
-1	-1	1


Sobel (5 x 5):

-1	-2	0	2	1
-2	-3	0	3	2
-3	-5	0	5	3
-2	-3	0	3	2
-1	-2	0	2	1

1	2	3	2	1
2	3	5	3	2
0	0	0	0	0
-2	-3	-5	-3	-2
-1	-2	-3	-2	-1



Poor Localization
Less Noise Sensitive
Good Detection



Gradients -> edges

Primary edge detection steps:

1. Smoothing: suppress noise
2. Edge enhancement: filter for contrast
3. Edge localization



Determine which local maxima from filter output are actually edges vs. noise

Thresholding and thinning



Thresholding

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)





Original image





Gradient magnitude image





Lower threshold





Higher threshold



Example: Laplacian of Gaussian (LoG) and Canny Edge Detector

Marr and Hildreth Filtering, 1980.

- Smooth Image with Gaussian Filter
- Applying the Laplacian for a Gaussian-filtered image can be done in one step of convolution.
- Find zero-crossings
- Find slope of zero-crossings
- Apply threshold to slope and mark edges

J. Canny. 1986

- Smooth Image with Gaussian filter
- Compute Derivative of filtered image
- Find Magnitude and Orientation of gradient
- Apply Non-max suppression
- Apply Thresholding (Hysteresis)



Canny edge detector



- This is probably the most widely used edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that optimizes the product of *signal-to-noise ratio* and localization

J. Canny, *A Computational Approach To Edge Detection*, IEEE Trans.
Pattern Analysis and Machine Intelligence, 8:679-714, 1986.



Canny edge detector



- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`



The Canny edge detector



original image (Lena)



The Canny edge detector



norm of the gradient






The Canny edge detector



thresholding





The Canny edge detector



How to turn
these thick
regions of the
gradient into
curves?



Non-maximum suppression



Check if pixel is local maximum along gradient direction,
select single max across width of the edge



- requires checking interpolated pixels p and r



Bilinear Interpolation



$$f(x, y) \approx [1 - x \quad x] \begin{bmatrix} f(0, 0) & f(0, 1) \\ f(1, 0) & f(1, 1) \end{bmatrix} \begin{bmatrix} 1 - y \\ y \end{bmatrix}.$$






Sidebar: Interpolation options



- `imx2 = imresize(im, 2, interpolation_type)`
- ‘nearest’
 - Copy value from nearest known
 - Very fast but creates blocky edges
- ‘bilinear’
 - Weighted average from four nearest known pixels
 - Fast and reasonable results
- ‘bicubic’ (default)
 - Non-linear smoothing over larger area
 - Slower, visually appealing, may create negative pixel values



Examples from http://en.wikipedia.org/wiki/Bicubic_interpolation



Before non-max suppression





After non-max suppression



Problem:
pixels along
this edge
didn't
survive the
thresholding



Hysteresis thresholding



- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels






Closing edge gaps




- Check that maximum value of gradient value is sufficiently large
 - drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.





Hysteresis thresholding



Source: S. Seitz



Hysteresis thresholding



original image



high threshold
(strong edges)



low threshold
(weak edges)



hysteresis threshold



Final Canny Edges





Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$




Canny with $\sigma = 2$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features




Example: Canny Edge Detection





Example: Canny Edge Detection





Recap: Canny edge detector



- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')` ;
- `>>help edge`



Low-level edges vs. perceived contours



Background

Texture

Shadows₉₀






Low-level edges vs. perceived contours





image



human segmentation



gradient magnitude










- Berkeley segmentation database:




<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

Protocol

You will be presented a photographic image. Divide the image into some number of segments, where the segments represent “things” or “parts of things” in the scene. The number of segments is up to you, as it depends on the image. Something between 2 and 30 is likely to be appropriate. It is important that all of the segments have approximately equal importance.


- Custom segmentation tool
- Subjects obtained from work-study program
(UC Berkeley undergraduates)






Consistency

A



Perceptual organization forms a tree:



B


C

- A,C are refinements of B
- A,C are mutual refinements
- A,B,C represent the same percept
 - Attention accounts for differences


★ Two segmentations are consistent when they can be explained by the same segmentation tree (i.e. they could be derived from a single perceptual organization).



Learn from humans which combination of features is most indicative of a “good” contour?



Contours can be defined by any of a number of cues (P. Cavanagh)





Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues

David R. Martin, *Member, IEEE*, Charless C. Fowlkes, and Jitendra Malik, *Member, IEEE*

Abstract—The goal of this work is to accurately detect and localize boundaries in natural scenes using local image measurements. We formulate features that respond to characteristic changes in brightness, color, and texture associated with natural boundaries. In order to combine the information from these features in an optimal way, we train a classifier using human labeled images as ground truth. The output of this classifier provides the posterior probability of a boundary at each image location and orientation. We present precision-recall curves showing that the resulting detector significantly outperforms existing approaches. Our two main results are 1) that cue combination can be performed adequately with a simple linear model and 2) that a proper, explicit treatment of texture is required to detect boundaries in natural images.

Index Terms—Texture, supervised learning, cue combination, natural images, ground truth segmentation data set, boundary detection, boundary localization.

Dataflow

Image



Boundary Cues

Brightness


Color

Texture

Cue Combination

Model

P_b




Challenges: texture cue, cue combination

Goal: learn the posterior probability of a boundary
 $P_b(x,y,\theta)$ from local information only




What features are responsible for perceived edges?





What features are responsible for perceived edges?



Feature profiles (oriented energy, brightness, color, and texture gradients) along the patch's horizontal diameter

Dataflow

Image



Optimized Cues

Brightness


Color

Texture


Cue Combination

Model


P_b



Human Segmentations





Oriented Feature Gradient




Individual Features

- 1976 CIE L*a*b* colorspace
- Brightness Gradient $BG(x,y,r,\theta)$
 - Difference of L* distributions
- Color Gradient $CG(x,y,r,\theta)$
 - Difference of a*b* distributions
- Texture Gradient $TG(x,y,r,\theta)$
 - Difference of distributions of V1-like filter responses



These are combined using logistic regression


Filter Outputs



2D Textons

- Goal: find canonical local features in a texture;

- 1) Filter image with linear filters:




- 2) Vector quantization (k-means) on filter outputs;

- 3) Quantization centers are the textons.


- Spatial distribution of textons defines the texture;

Texture gradient = Chi square distance between
texton histograms in half disks across edge



Various Cue Combinations


(Martin, Fowlkes,
Malik, 2004)



Image



BG+CCG+TG




Human





Brightness



Combined



Human



Results




Human (0.95)








Results



Human (0.96)







For more:
<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>



Contour Detection and Hierarchical Image Segmentation

Pablo Arbeláez, *Member, IEEE*, Michael Maire, *Member, IEEE*,
Charless Fowlkes, *Member, IEEE*, and Jitendra Malik, *Fellow, IEEE*.

Abstract—This paper investigates two fundamental problems in computer vision: contour detection and image segmentation. We present state-of-the-art algorithms for both of these tasks. Our contour detector combines multiple local cues into a globalization framework based on spectral clustering. Our segmentation algorithm consists of generic machinery for transforming the output of any contour detector into a hierarchical region tree. In this manner, we reduce the problem of image segmentation to that of contour detection. Extensive experimental evaluation demonstrates that both our contour detection and segmentation methods significantly outperform competing algorithms. The automatically generated hierarchical segmentations can be interactively refined by user-specified annotations. Computation at multiple image resolutions provides a means of coupling our system to recognition applications.



Spectral Pb





Fig. 7. **Spectral Pb.** **Left:** Image. **Middle Left:** The thinned non-max suppressed multiscale Pb signal defines a sparse affinity matrix connecting pixels within a fixed radius. Pixels i and j have a low affinity as a strong boundary separates them, whereas i and k have high affinity. **Middle:** First four generalized eigenvectors resulting from spectral clustering. **Middle Right:** Partitioning the image by running K-means clustering on the eigenvectors erroneously breaks smooth regions. **Right:** Instead, we compute gradients of the eigenvectors, transforming them back into a contour signal.




Global pB boundary detector




Extract Pb



Compute Eigenvectors



Gradient of
eigenvectors







Global pB boundary detector



Fig. 9. **Benefits of globalization.** When compared with the local detector P_b , our detector gP_b reduces clutter and completes contours. The thresholds shown correspond to the points of maximal F-measure on the curves in Figure 1.







Fast Edge Detection Using Structured Forests

Piotr Dollár and C. Lawrence Zitnick

Abstract—Edge detection is a critical component of many vision systems, including object detectors and image segmentation algorithms. Patches of edges exhibit well-known forms of local structure, such as straight lines or T-junctions. In this paper we take advantage of the structure present in local image patches to learn both an accurate and computationally efficient edge detector. We formulate the problem of predicting local edge masks in a structured learning framework applied to random decision forests. Our novel approach to learning decision trees robustly maps the structured labels to a discrete space on which standard information gain measures may be evaluated. The result is an approach that obtains realtime performance that is orders of magnitude faster than many competing state-of-the-art approaches, while also achieving state-of-the-art edge detection results on the BSDS500 Segmentation dataset and NYU Depth dataset. Finally, we show the potential of our approach as a general purpose edge detector by showing our learned edge models generalize well across datasets.

Index Terms—Edge detection, segmentation, structured random forests, real-time systems, visual features








Fig. 1. Edge detection results using three versions of our Structured Edge (SE) detector demonstrating tradeoffs in accuracy vs. runtime. We obtain realtime performance while simultaneously achieving state-of-the-art results. ODS numbers were computed on BSDS [1] on which the popular gPb detector [1] achieves a score of .73. The variants shown include SE, SE+SH, and SE+MS+SH, see §4 for details.



Edge Detection with Structured Random Forests



- Goal: quickly predict whether each pixel is an edge
- Insights
 - Predictions can be learned from training data
 - Predictions for nearby pixels should not be independent
- Solution
 - Train structured random forests to split data into patches with similar boundaries based on features
 - Predict boundaries at patch level, rather than pixel level, and aggregate (average votes)



Boundaries
in patch


124



Edge Detection with Structured Random Forests

Algorithm

1. Extract overlapping 32x32 patches at three scales
2. Features are pixel values and pairwise differences in feature maps (LUV color, gradient magnitude, oriented gradient)
3. Predict T boundary maps in the central 16x16 region using T trained decision trees
4. Average predictions for each pixel across all patches







Fig. 2. Illustration of the decision tree node splits: (a) Given a set of structured labels such as segments, a splitting function must be determined. Intuitively a good split (b) groups similar segments, whereas a bad split (c) does not. In practice we cluster the structured labels into two classes (d). Given the class labels, a standard splitting criterion, such as Gini impurity, may be used (e).





Edge Detection with Structured Random Forests



Results

BSDS 500

	ODS	OIS	AP	FPS
Human	.80	.80	-	-
Canny	.60	.64	.58	15
Felz-Hutt [11]	.61	.64	.56	10
Hidayat-Green [16]	.62 [†]	-	-	20
BEL [9]	.66 [†]	-	-	1/10
gPb + GPU [6]	.70 [†]	-	-	1/2 [‡]
gPb [1]	.71	.74	.65	1/240
gPb-owt-ucm [1]	.73	.76	.73	1/240
Sketch tokens [21]	.73	.75	.78	1
SCG [31]	.74	.76	.77	1/280
SE-SS, $T=1$.72	.74	.77	60
SE-SS, $T=4$.73	.75	.77	30
SE-MS, $T=4$.74	.76	.78	6

NYU Depth dataset edges

	ODS	OIS	AP	FPS
gPb [1] (rgb)	.51	.52	.37	1/240
SCG [31] (rgb)	.55	.57	.46	1/280
SE-SS (rgb)	.58	.59	.53	30
SE-MS (rgb)	.60	.61	.56	6
gPb [1] (depth)	.44	.46	.28	1/240
SCG [31] (depth)	.53	.54	.45	1/280
SE-SS (depth)	.57	.58	.54	30
SE-MS (depth)	.58	.59	.57	6
gPb [1] (rgbd)	.53	.54	.40	1/240
SCG [31] (rgbd)	.62	.63	.54	1/280
SE-SS (rgbd)	.62	.63	.59	25
SE-MS (rgbd)	.64	.65	.63	5



Holistically-Nested Edge Detection

Saining Xie

Dept. of CSE and Dept. of CogSci
University of California, San Diego
9500 Gilman Drive, La Jolla, CA 92093

s9xie@eng.ucsd.edu

Zhuowen Tu

Dept. of CogSci and Dept. of CSE
University of California, San Diego
9500 Gilman Drive, La Jolla, CA 92093

ztu@ucsd.edu




Figure 1. Illustration of the proposed HED algorithm. In the first row: (a) shows an example test image in the BSD500 dataset [28]; (b) shows its corresponding edges as annotated by human subjects; (c) displays the HED results. In the second row: (d), (e), and (f), respectively, show side edge responses from layers 2, 3, and 4 of our convolutional neural networks. In the third row: (g), (h), and (i), respectively, show edge responses from the Canny detector [4] at the scales $\sigma = 2.0$, $\sigma = 4.0$, and $\sigma = 8.0$. HED shows a clear advantage in consistency over Canny.



Holistically nested edge detection





Table 4. Results on BSDS500. *BSDS300 results, †GPU time

	ODS	OIS	AP	FPS
Human	.80	.80	-	-
Canny	.600	.640	.580	15
Felz-Hutt [9]	.610	.640	.560	10
BEL [5]	.660*	-	-	1/10
gPb-owt-ucm [1]	.726	.757	.696	1/240
Sketch Tokens [24]	.727	.746	.780	1
SCG [31]	.739	.758	.773	1/280
SE-Var [6]	.746	.767	.803	2.5
OEF [13]	.749	.772	.817	-
DeepNets [21]	.738	.759	.758	1/5†
N4-Fields [10]	.753	.769	.784	1/6†
DeepEdge [2]	.753	.772	.807	1/10 ³ †
CSCNN [19]	.756	.775	.798	-
DeepContour [34]	.756	.773	.797	1/30†
HED (ours)	.782	.804	.833	2.5†, 1/12





State of edge detection



- Local edge detection is mostly solved
 - Intensity gradient, color, texture
 - HED on BSDS 500 is near human performance
- Some room for improvement by taking advantage of higher-level knowledge (e.g., objects)
- Still hard to produce all objects within a small number of regions



Finding straight lines





Finding line segments using connected components



1. Compute canny edges
 - Compute: gx, gy (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(gy / gx)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x, y 2nd moment matrix of their points

$$\mathbf{M} = \begin{bmatrix} \sum(x - \mu_x)^2 & \sum(x - \mu_x)(y - \mu_y) \\ \sum(x - \mu_x)(y - \mu_y) & \sum(y - \mu_y)^2 \end{bmatrix} \quad [v, \lambda] = \text{eig}(\mathbf{M})$$

Larger eigenvector
↓

$$\theta = \text{atan} 2(v(2,2), v(1,2))$$
$$conf = \lambda_2 / \lambda_1$$

5. Threshold on straightness, store segment



Things to remember

- Canny edge detector = smooth → derivative → thin → threshold → link
- Pb: learns weighting of gradient, color, texture differences
 - More recent learning approaches give at least as good accuracy and are faster
- Straight line detector = canny + gradient orientations → orientation binning → linking → check for straightness

